Comp 251: Assignment 2

Answers must be submitted online by November 13th (11:55:00 pm), 2020.

General instructions (Read carefully!)

- **Important:** All of the work you submit must be done by only you, and your work must not be submitted by someone else. Plagiarism is academic fraud and is taken very seriously. For Comp251, we will use software that compares programs for evidence of similar code. This software is very effective and it is able to identify similarities in the code even if you change the name of your variables and the position of your functions. The time that you will spend modifying your code, would be better invested in creating an original solution.

  Please don't copy. We want you to succeed and are here to help. Here are a couple of general guidelines to help you avoid plagiarism:

  Never look at another assignment solution, whether it is on paper or on the computer screen. Never share your assignment solution with another student. This applies to all drafts of a solution and to incomplete solutions. If you find code on the web, or get code from a private tutor, that solves part or all of an assignment, do not use or submit any part of it! A large percentage of the academic offenses in CS involve students who have never met, and who just happened to find the same solution online, or work with the same tutor. If you find a solution, someone else will too. The easiest way to avoid plagiarism is to only discuss a piece of work with the Comp251 TAs, the CS Help Centre TAs, or the COMP 251 instructors.

- Your solution must be submitted electronically on codePost. Here is a short **tutorial** to help you understand how the platform works. You will receive an email inviting you to join the class there in early October.

- To some extent, collaborations are allowed. These collaborations should not go as far as sharing code or giving away the answer. You must indicate on your assignments (i.e. as a comment at the beginning of your java source file) the names of the people with whom you collaborated or discussed your assignments (including members of the course staff). If you did not collaborate with anyone, you write "No collaborators". If asked, you should be able to orally explain your solution to a member of the course staff.

- This assignment is due on October 16^{th} at 11h59:59 pm. It is your responsibility to guarantee that your assignment is submitted on time. We do not cover technical issues or unexpected difficulties you may encounter. Last minute submissions are at your own risk.

- This assignment includes a programming component, which counts for 100% of the grade, and an optional long answer component designed to prepare you for the exams. This component will not be graded, but a solution guide will be published.

- Multiple submissions are allowed before the deadline. We will only grade the last submitted file. Therefore, we encourage you to submit as early as possible a preliminary version of your solution to avoid any last minute issue.

- Late submissions can be submitted for 24 hours after the deadline, and will receive a flat penalty of 20%. We will not accept any submission more than 24 hours after the deadline. The submission site will be closed, and there will be no exceptions, except medical.

- In exceptional circumstances, we can grant a small extension of the deadline (e.g. 24h) for medical reasons only. However, such request must be submitted before the deadline, and justified by a medical note from a doctor, which must also be submitted to the McGill administration.

- Violation of any of the rules above may result in penalties or even absence of grading. If anything is unclear, it is up to you to clarify it by asking either directly the course staff during office hours, by email at (cs251@cs.mcgill.ca) or on the discussion board on Piazza (recommended). Please, note that we reserve the right to make specific/targeted announcements affecting/extending these rules in class and/or on the website. It is your responsibility to monitor Piazza for announcements.

- The course staff will answer questions about the assignment during office hours or in the online forum. We urge you to ask your questions as early as possible. We cannot guarantee that questions asked less than 24h before the submission deadline will be answered in time. In particular, we will not answer individual emails about the assignment that are sent sent the day of the deadline.

Programming component

- You are provided some starter code that you should fill in as requested. Add your code only where you are instructed to do so. You can add some helper methods. Do not modify the code in any other way and in particular, do not change the methods or constructors that are already given to you, do not import extra code and do not touch the method headers. The format that you see on the provided code is the only format accepted for programming questions. **Any failure to comply with these rules will result in an automatic 0.**

- Public tests cases are available on codePost. You can run them on your code at any time. If your code fails those tests, it means that there is a mistake somewhere. Even if your code passes those tests, it may still contain some errors. We will grade your code with a more challenging, private set of test cases. We therefore highly encourage you to modify that tester class, expand it and share it with other students on the discussion board. Do not include it in your submission.

- Your code should be properly commented and indented.

- **Do not change or alter the name of the files you must submit, or the method headers in these files**. Files with the wrong name will not be graded. Make sure you are not changing file names by duplicating them. For example, main (2).java will not be graded. Make sure to double-check your zip file.

- **You can submit either a zip file or individual files on codePost.** If you get more than 0 on the public tests, it means codePost accepted your files.

- **You will automatically get 0 if the files you submitted on codePost do not compile, since you can ensure yourself that they do**. Note that public test cases do not cover every situation and your code may crash when tested on a method that is not checked by the public tests. This is why you need to add your own test cases and compile and run your code from command line on linux.

Exercise 1 (Disjoint sets (24 points)) We want to implement a disjoint set data structure with union and find operations. The template for this program is available on the course website and named `DisjointSets.java`.

In this question, we model a partition of n elements with distinct integers ranging from 0 to $n-1$ (i.e. each element is represented by an integer in $[0, \cdots, n-1]$, and each integer in $[0, \cdots, n-1]$ represent one element). We choose to represent the disjoint sets with trees, and to implement the forest of trees with an array named `par`. More precisely, the value stored in `par[i]` is parent of the element `i`, and `par[i]==i` when `i` is the root of the tree and thus the representative of the disjoint set.

You will implement union by rank and the *path compression* technique seen in class. The rank is an integer associated with each node. Initially (i.e. when the set contains one single object) its value is 0. Union operations link the root of the tree with smaller rank to the root of the tree with larger rank. In the case where the rank of both trees is the same, the rank of the new root increases by 1. You can implement the rank with a specific array (called `rank`) that has been added to the template, or use the array `par` (this is tricky). Note that path compression does not change the rank of a node.

Download the file `DisjointSets.java`, and complete the methods `find(int i)` as well as `union(int i, int j)`. The constructor takes one argument n (a strictly positive integer) that indicates the number of elements in the partition, and initializes it by assigning a separate set to each element.

The method `find(int i)` will return the representative of the disjoint set that contains i (do not forget to implement path compression here!). The method `union(int i, int j)` will merge the set with smaller rank (for instance `i`) in the disjoint set with larger rank (in that case `j`). In that case, the root of the tree containing `i` will become a child of the root of the tree containing `j`), and return the representative (as an integer) of the new merged set. Do not forget to update the ranks. In the case where the ranks are identical, you will merge `i` into `j`.

Once completed, compile and run the file `DisjointSets.java`. It should produce the output available in the file `unionfind.txt` available on the course website.
Note: You will need to complete this question to implement Exercise 2.

Exercise 2 (Minimum Spanning trees (24 points)) We will implement the Kruskal algorithm to calculate the minimum spanning tree (MST) of an undirected weighted graph. Here you will use the file `DisjointSets.java` completed in the previous question and two other files: `WGraph.java` and `Kruskal.java` available on the course website. You will need the classes `DisjointSets` and `WGraph` to execute `Kruskal.java`. Your role will be to complete two methods in the template `Kruskal.java`.

The file `WGraph.java` implements two classes `WGraph` and `Edge`. An `Edge` object stores all informations about edges (i.e. the two vertices and the weight of the edge), which are used to build graphs.

The class `WGraph` has two constructors `WGraph()` and `WGraph(String file)`. The first one creates an empty graph and the second uses a file to initialize a graph. Graphs are encoded using the following format. The first line of this file is a single integer n that indicates the number of nodes in the graph. Each vertex is labelled with an number in $[0, \cdots, n-1]$, and each integer in $[0, \cdots, n-1]$ represents one and only one vertex. The following lines respect the syntax "n_1 n_2 w", where n_1 and n_2 are integers representing the nodes connected by an edge, and w the weight of this edge. n_1, n_2, and w must be separated by space(s). An example of such file can be found on the course website with the file `g1.txt`. These files will be used as an input in the program `Kruskal.java` to initialize the graphs. Thus, an example of a command line is `java Kruskal g1.txt`.

The class `WGraph` also provide a method `addEdge(Edge e)` that adds an edge to a graph (i.e. an object of this class). Another method `listOfEdgesSorted()` allows you to access the list of edges

of a graph in the increasing order of their weight.

You task will be to complete the two static methods `isSafe(DisjointSets p, Edge e)` and `kruskal(WGraph g)` in `Kruskal.java`. The method `isSafe` considers a partition of the nodes p and an edge e, and should return `True` if it is safe to add the edge e to the MST, and `False` otherwise. The method `kruskal` will take a graph object of the class `WGraph` as an input, and return another `WGraph` object which will be the MST of the input graph.

Once completed, compile all the java files and run the command line `java Kruskal g1.txt`. An example of the expected output is available in the file `mst1.txt`. You are invited to run other examples of your own to verify that your program is correct. Note that you need to compile it with the two other files for it to work.

Exercise 3 (Greedy algorithms (52 points)) In this exercise, you will plan your homework with a greedy algorithm. The input is a list of homeworks defined by two arrays: deadlines and weights (the relative importance of the homework towards your final grade). These arrays have the same size and they contain integers between 1 and 100. The index of each entry in the arrays represents a single homework, for example, `Homework 2` is defined as a homework with deadline `deadlines[2]` and weight `weights[2]`. Each homework takes exactly one hour to complete.

Your task is to output a `homeworkPlan`: an array of length equal to the last deadline. Each entry in the array represents a one-hour timeslot, starting at 0 and ending at 'last deadline - 1'. For each time slot, `homeworkPlan` indicates the homework which you plan to do during that slot. You can only complete a single homework in one 1-hour slot. The homeworks are due at the beginning of a time slot, in other words if an assignment's deadline is x, then the last time slot when you can do it is x - 1. For example, if the homework is due at t=14, then you can complete it before or during the slot t=13. If your solution plans to do `Homework 2` first, then you should have `homeworkPlan[0]=2` in the output. Note that sometimes you will be given too much homework to complete in time, and that is okay.

Your homework plan should maximize the sum of the weights of completed assignments.

To organize your schedule, we give you a class `HW_Sched.java`, which defines an `Assignment` object, with a number (its index in the input array), a weight and a deadline.

The input arrays are unsorted. As part of the greedy algorithm, the template we provide sorts the homeworks using Java's `Collections.sort()`. This sort function uses Java's `compare()` method, which takes two objects as input, compares them, and outputs the order they should appear in. The template will ask you to override this `compare()` method, which will alter the way in which Assignments will be ordered. You have to determine what comparison criterion you want to use. Given two assignments A1 and A2, the method should output:

- 0, if the two items are equivalent

- 1, if a1 should appear after a2 in the sorted list

- -1, if a2 should appear after a1 in the sorted list

The `compare` method (called by `Collections.sort()`) should be the only tool you use to re-organize lists and arrays in this problem. You will then implement the rest of the `SelectAssignments()` method.

Submit all the Java files you modified on codePost.

Exercise 4 (Change-Making Problem (0 points) In this problem, we give a formal look at the problem of returning a given amount of money using the minimum number of coins (including bills), for a given coin system.

For example, the best way to give someone 7 dollars is to add up a 5 dollars bill and a 2 dollars coin. For simplicity, we will consider all denominations to be coins.

First, let us define the problem: a coin system is a m-tuple $c = (c_1, c_2, \ldots, c_m)$ such that $c_1 > c_2 > \cdots > c_m = 1$. For a given coin system c and a positive integer x representing the amount of money we want to gather, we want to find a solution (a m-tuple of non-negative integers) $k = (k_1, k_2, \ldots, k_m)$ such that $x = \sum_{i=1}^{m} k_i c_i$ so as to minimize $\sum_{i=1}^{m} k_i$.

There exists a greedy algorithm to find the optimal solution for certain coin systems: for a given x, we select the largest coin $c_i \leq x$. Then, we repeat this step for $x - c_i$, and continue repeating it until x becomes 0. For instance, with the coin system $(10, 5, 2, 1)$, the algorithm decomposes $x = 27$ into $10, 10, 5$, and 2.

We describe a coin system as *canonical* if and only if the solution given by the greedy algorithm described above is optimal for any positive integer x. For example, all systems $(a, 1)$ with $a > 1$ are canonical. For any positive integer x, we can express such a change system in the form of Euclidean division: $x = aq + r$ with $r < a$. The greedy solution for x is then the tuple $g = (q, r)$. To prove that the solution is optimal, we can proceed as follows:

Let's consider $g' = (q', r')$ different than g such that $x = aq' + r'$. Because q is already at its highest value under x and cannot be above x, we have $q' < q$, otherwise $q' = q$ causes $g' = g$. Since $(q' + r') - (q + r) = (q' + x - aq') - (q + x - aq) = (a - 1)(q - q') > 0$, g' would sum up to more coins than the initial solution g, for any g' satisfying the problem definition. Thus, the solution g is optimal, and the system $(a, 1)$ is canonical.

4.1 (0 points) Design a non-canonical system of 3-tuple $c = (c_1, c_2, c_3)$ and prove your system is non-canonical.

4.2 (0 points) Let q and n be two integers ≥ 2. Prove that the system $c = (q^n, q^{n-1}, \ldots, q, 1)$ is canonical.

4.3 (0 bonus points) Prove that the Euro system $c = (200, 100, 50, 20, 10, 5, 2, 1)$ is canonical. There will be no partial credit for this question.